

COMP 4511 - System and Kernel Programming in Linux Programming Assignment #3

Assignment 3 - Simple Linux Shell (Part 2)

You should modify the simple Linux Shell (Part 1) with the addition of 2 key features:

- I/O redirection
- shell pipes

Please note that you should not break existing features in your submitted part 1.

Feature 1 - I/O Redirection

To accomplish this feature, you need to use the `dup()` or `dup2()` (recommended).

A process has 3 default file identifiers:

- `stdin` (standard input)
- `stdout` (standard output)
- and `stderr` (standard error output)

If the process reads from `stdin`, then the data that it receives will be directed from the keyboard to the `stdin` file descriptor. Similarly, data received from `stdout` and `stderr` are mapped to the terminal display. In a typical shell program, a user can redefine `stdin` or `stdout` to another input/output source. A “less than” character (i.e. ‘<’) followed by a filename is used for input redirection. Similarly, a “greater than” character (i.e. ‘>’) followed by a filename is used for output redirection.

For example, a command such as:

```
[DIR] myshell> test < main.cpp > output.txt
```

This command will create a child process to execute the “test” command. Before the “test” command is launched, `stdin` will be redirected to read the input stream from “main.cpp” and `stdout` will be redirected to write the output stream to “output.txt”.

The shell can redirect I/O by manipulating the child process’s file descriptors. A newly created child process inherits the open file descriptors of its parent, specifically the same keyboard for `stdin` and the terminal display for `stdout` and `stderr`.

The shell can change the child's file descriptors so that it reads and writes to files rather than to the keyboard and display.

Feature 2 - Shell Pipes

A pipe is represented in the kernel by a file descriptor, just like a file. A process creates a pipe by:

```
int ps[2];  
pipe(ps);
```

Here, `ps[0]` is a file pointer (an index into the process's open file table) to the read end of the pipe and `ps[1]` is the file pointer to the write end of the pipe. For two or more processes to use pipes for IPC, a common ancestor of the process must create the pipe prior to creating the processes.

In a shell, a pipe symbol '|' is used to connect the output of the first command to the input of the second command. For example,

```
[DIR] myshell> ls | sort
```

The "ls" command lists the contents of the current directory. As the output of "ls" is already connected to "sort", it won't print out the content to the screen. After the output of ls has been sorted by sort, the sorted list of files appears on the screen.

Please note that you are required to support multiple pipes. For example:

```
[DIR] myshell> ps aux | grep root | sort | less
```

Sample output for I/O redirection and pipes

Sample output #1

```
[myshell] myshell> ./test < input.txt > out.txt  
[myshell] myshell> cat out.txt  
Hello World!  
Hello World!  
Hello World!  
Hello World!  
myshell> cat out.txt | wc -l  
4  
myshell>
```

The test program simply reads a file and outputs the file content to the screen. The `input.txt` includes 4 lines of "HelloWorld!"

Sample output #2

myshell should accept more than one pipes and support a combination of output redirection and pipes. In particular, myshell should support multiple pipes, as shown by the following example,

```
[myshell] myshell> ./test < input.txt | head -n 2 | wc -l >
output.txt
[myshell] myshell> nano output.txt
```

The content of output.txt now should be 2. Note that the output redirection is always associated with the command it follows and NOT with the whole command line.